📖 Review Convolutional Networks

Images are one of the most common modes of data. Making sense of images is one of the fundamental challenges of artificial intelligence. Popular applications include visual object recognition (e.g., of traffic signs in self-driving cars) or face recognition.

However, image data is very high dimensional. For example, a 1 megapixel image contains 3 million numbers. Images are usually represented as a 3 × height × width (or height × width × 3) tensor (or 3D matrices). The dimension with 3 elements is known as the color channel, and for a usual image we have 3 color channels for red, green, and blue. The other two dimensions capture the coordinates of each pixel in the image. For gray-scale images, we have a height × width or height × width × 1 image, since the color of the pixel can be captured by a single number describing its darkness.

If we have an image of size 3 × 1000 × 1000, then this image essentially has 3 million dimensions (3 × 1000 × 1000). A naive multilayer perception (i.e., plain neural network without convolutional layers) trained on such images will have many weights and nodes, so will need a huge number of photos to train it.

Not only are images high dimensional, but neighboring pixels in an image are also highly correlated. More concretely, an image pixel does not really contain a lot of information, but a block of image pixels is often much more meaningful. To understand this, consider an image of a tiger. Looking at a single pixel does not give you any information about the image, but by looking at a block of images (e.g., several neighboring pixels that make a patch of fur with orange and black stripes), you can guess that the image contains a tiger. If we vectorize an image before inputting it into a multilayer Perceptron, we end up with one long 1 million dimensional vector and essentially lose much of the spatial relation between neighboring pixels.

This is where convolutional neural networks (ConvNets) come into play. On a high level, ConvNets try to learn small filters that can detect small regional patterns at the first few layers. The vector output of these layers might have higher values if a particular pattern is present, and lower otherwise. In later layers, these outputs are passed into multilayer Perceptrons for classification.

In this module, we will explore different components of ConvNets and how we can piece them together to develop a special type of neural network that is fascinatingly well-suited for image data.

☆ **Key Points**

Images are represented as 3-dimensional tensors, of dimensions height × width × 3, where the last dimension captures the color channels.

Convolutional neural networks have the ability to filter local regions of images to recognize patterns.

◂ Previous

Next ▸